

Actor Model

#SCBCN14

sponsored by



Ignasi Marimon-Clos (@ignasi35)



thanks!

Questions!

about you



about me

@ignasi35

- 1) problem solver, Garbage Collector, mostly scala, java 8, agile for tech teams
- 2) kayaker
- 3) under construction
- 4) all things JVM

@ignasi35



Seven Concurrency Models in Seven Weeks

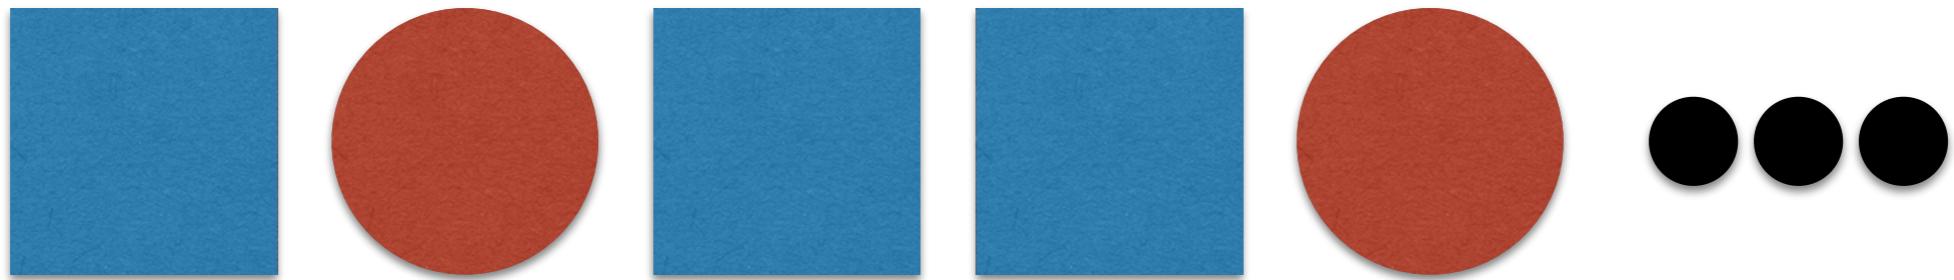
When Threads Unravel



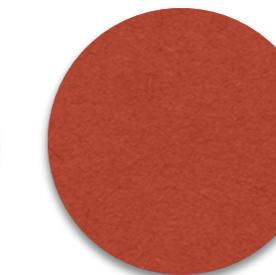
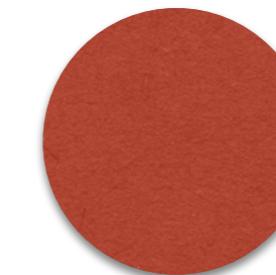
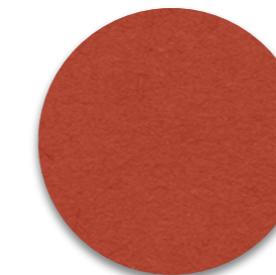
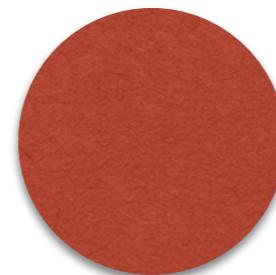
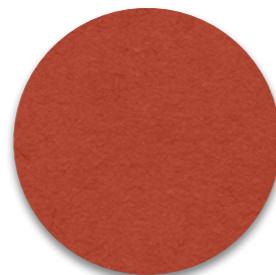
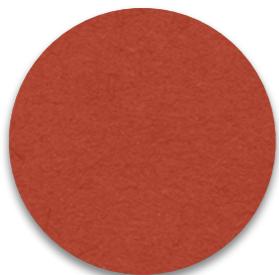
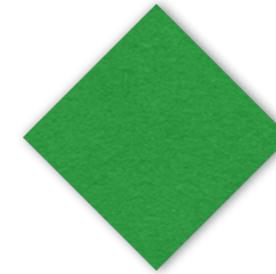
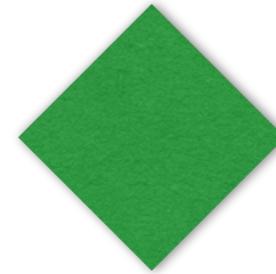
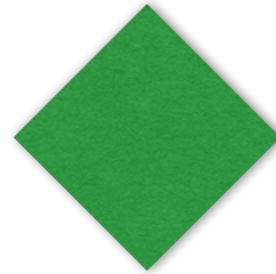
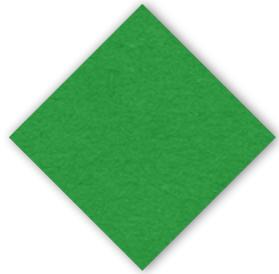
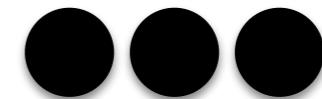
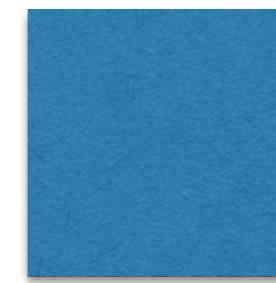
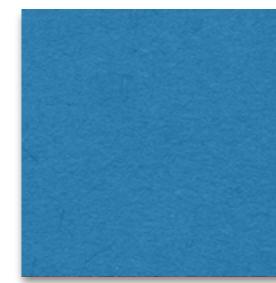
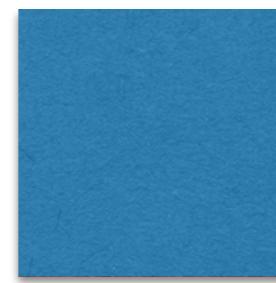
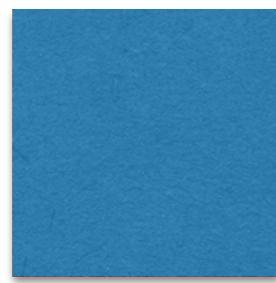
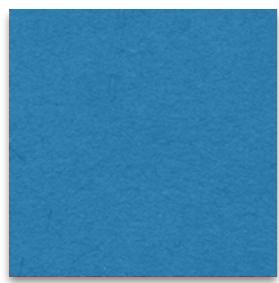
Paul Butcher

Series editor: *Bruce Tate*
Development editor: *Jacquelyn Carter*

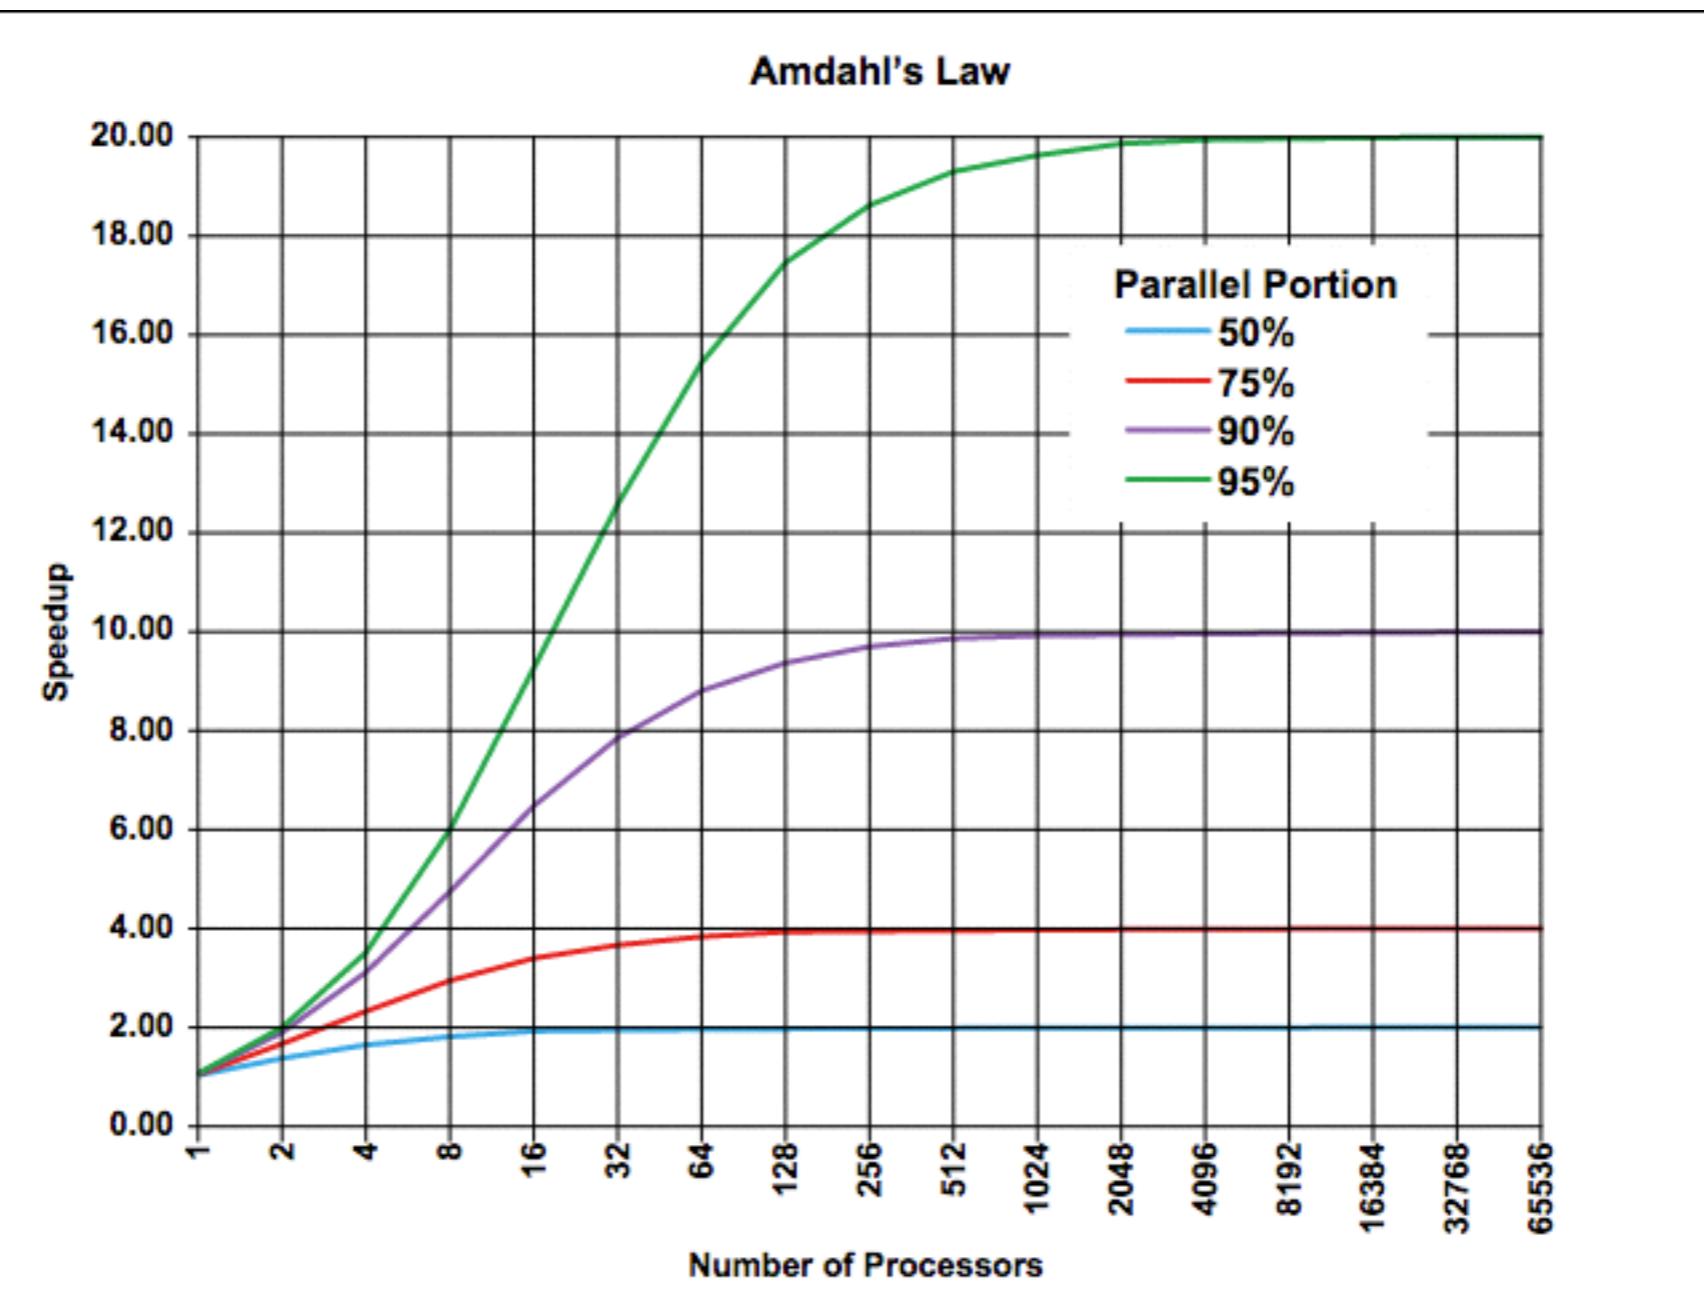
Concurrency



Parallelism



The Problem

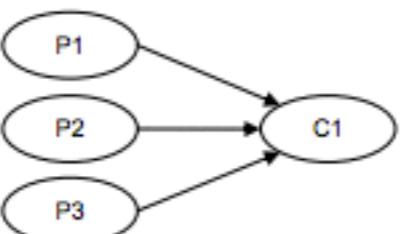




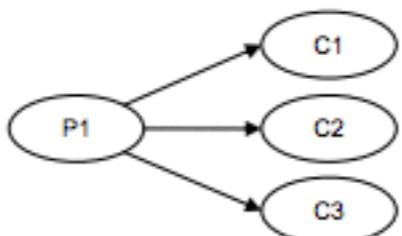
Unicast: 1P – 1C



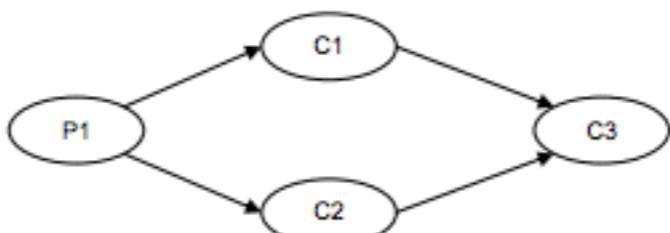
Three Step Pipeline: 1P – 3C



Sequencer: 3P – 1C

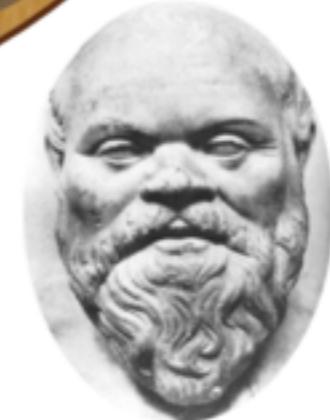
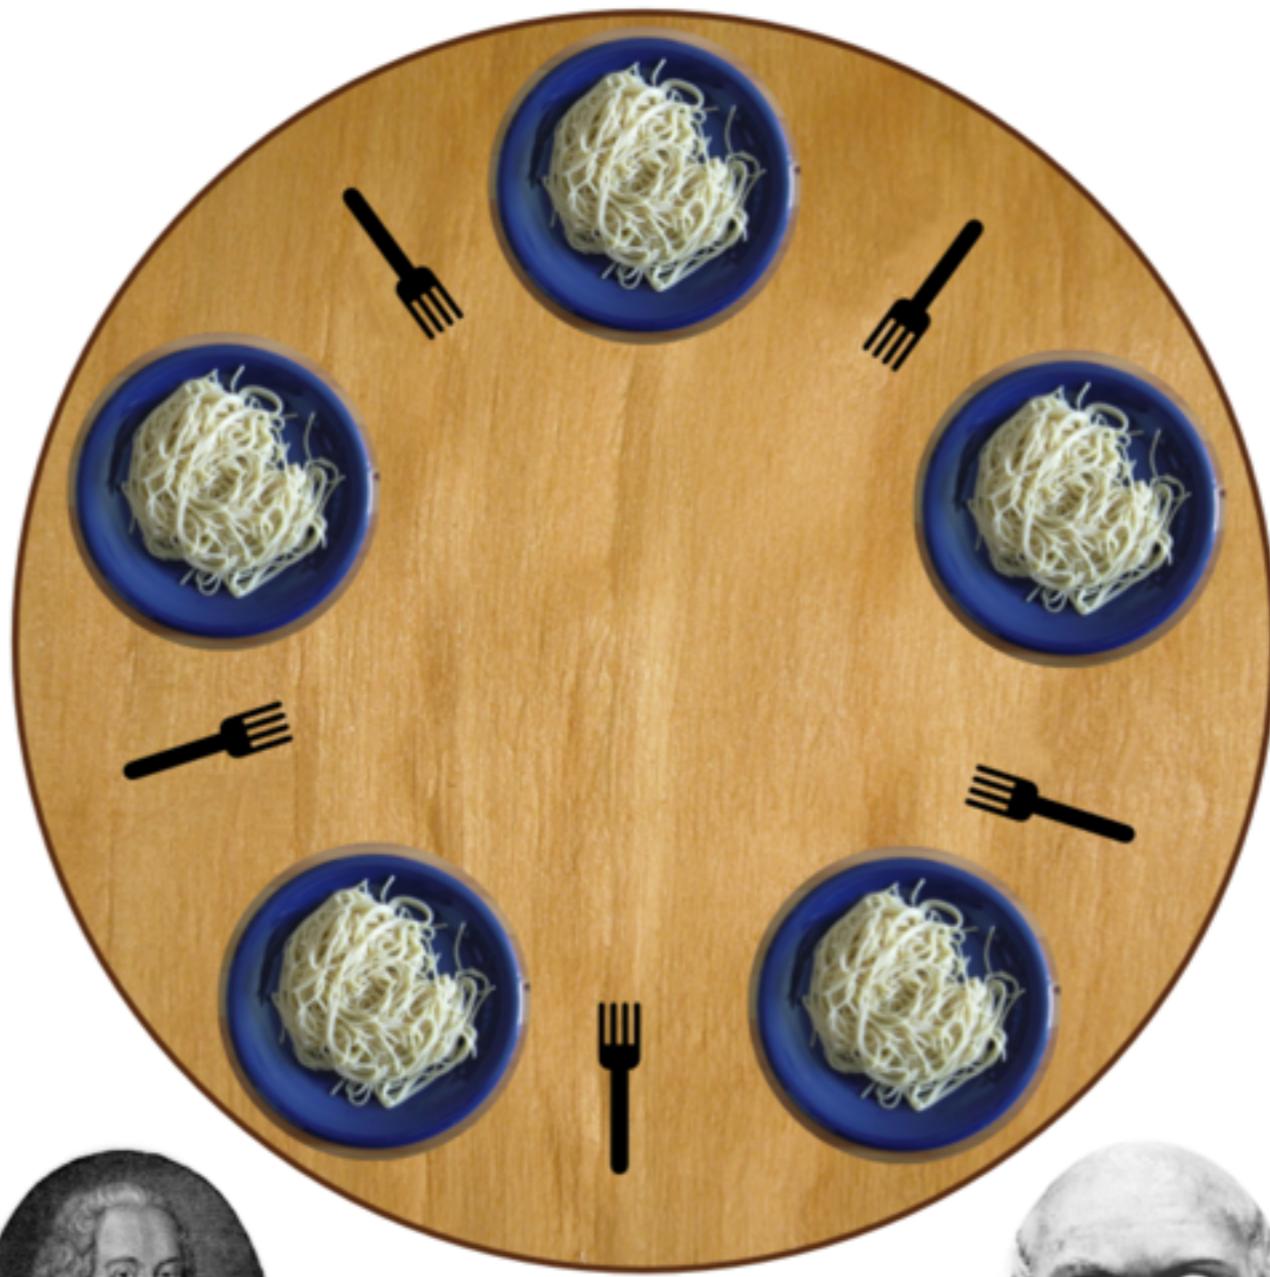


Multicast: 1P – 3C



Diamond: 1P – 3C

- CPU-bound or IO bound
- Latency or Throughput
 - memory footprint



@ignasi35

Threads and Locks

Threads and Locks

- low level
- debugging
- ThreadPool

Threads and Locks

- threads are heavy
- locks expensive
- low level
- no recovery
- locking granularity
- visibility

Threads and Locks

Mutual Exclusion
Lock Ordering
Memory Model



HD

Event Loop

Event Loop



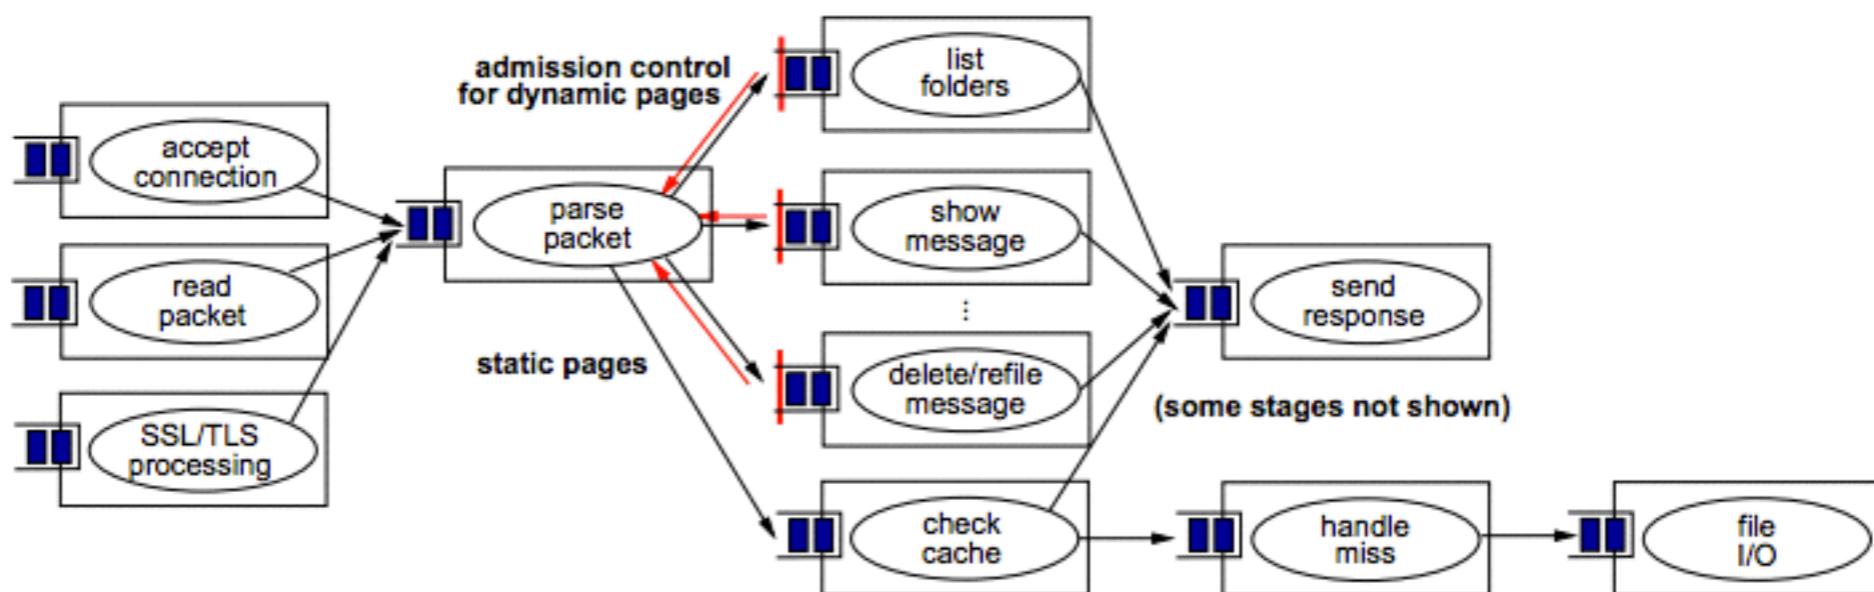


filmax

SEDA

@ignasi35

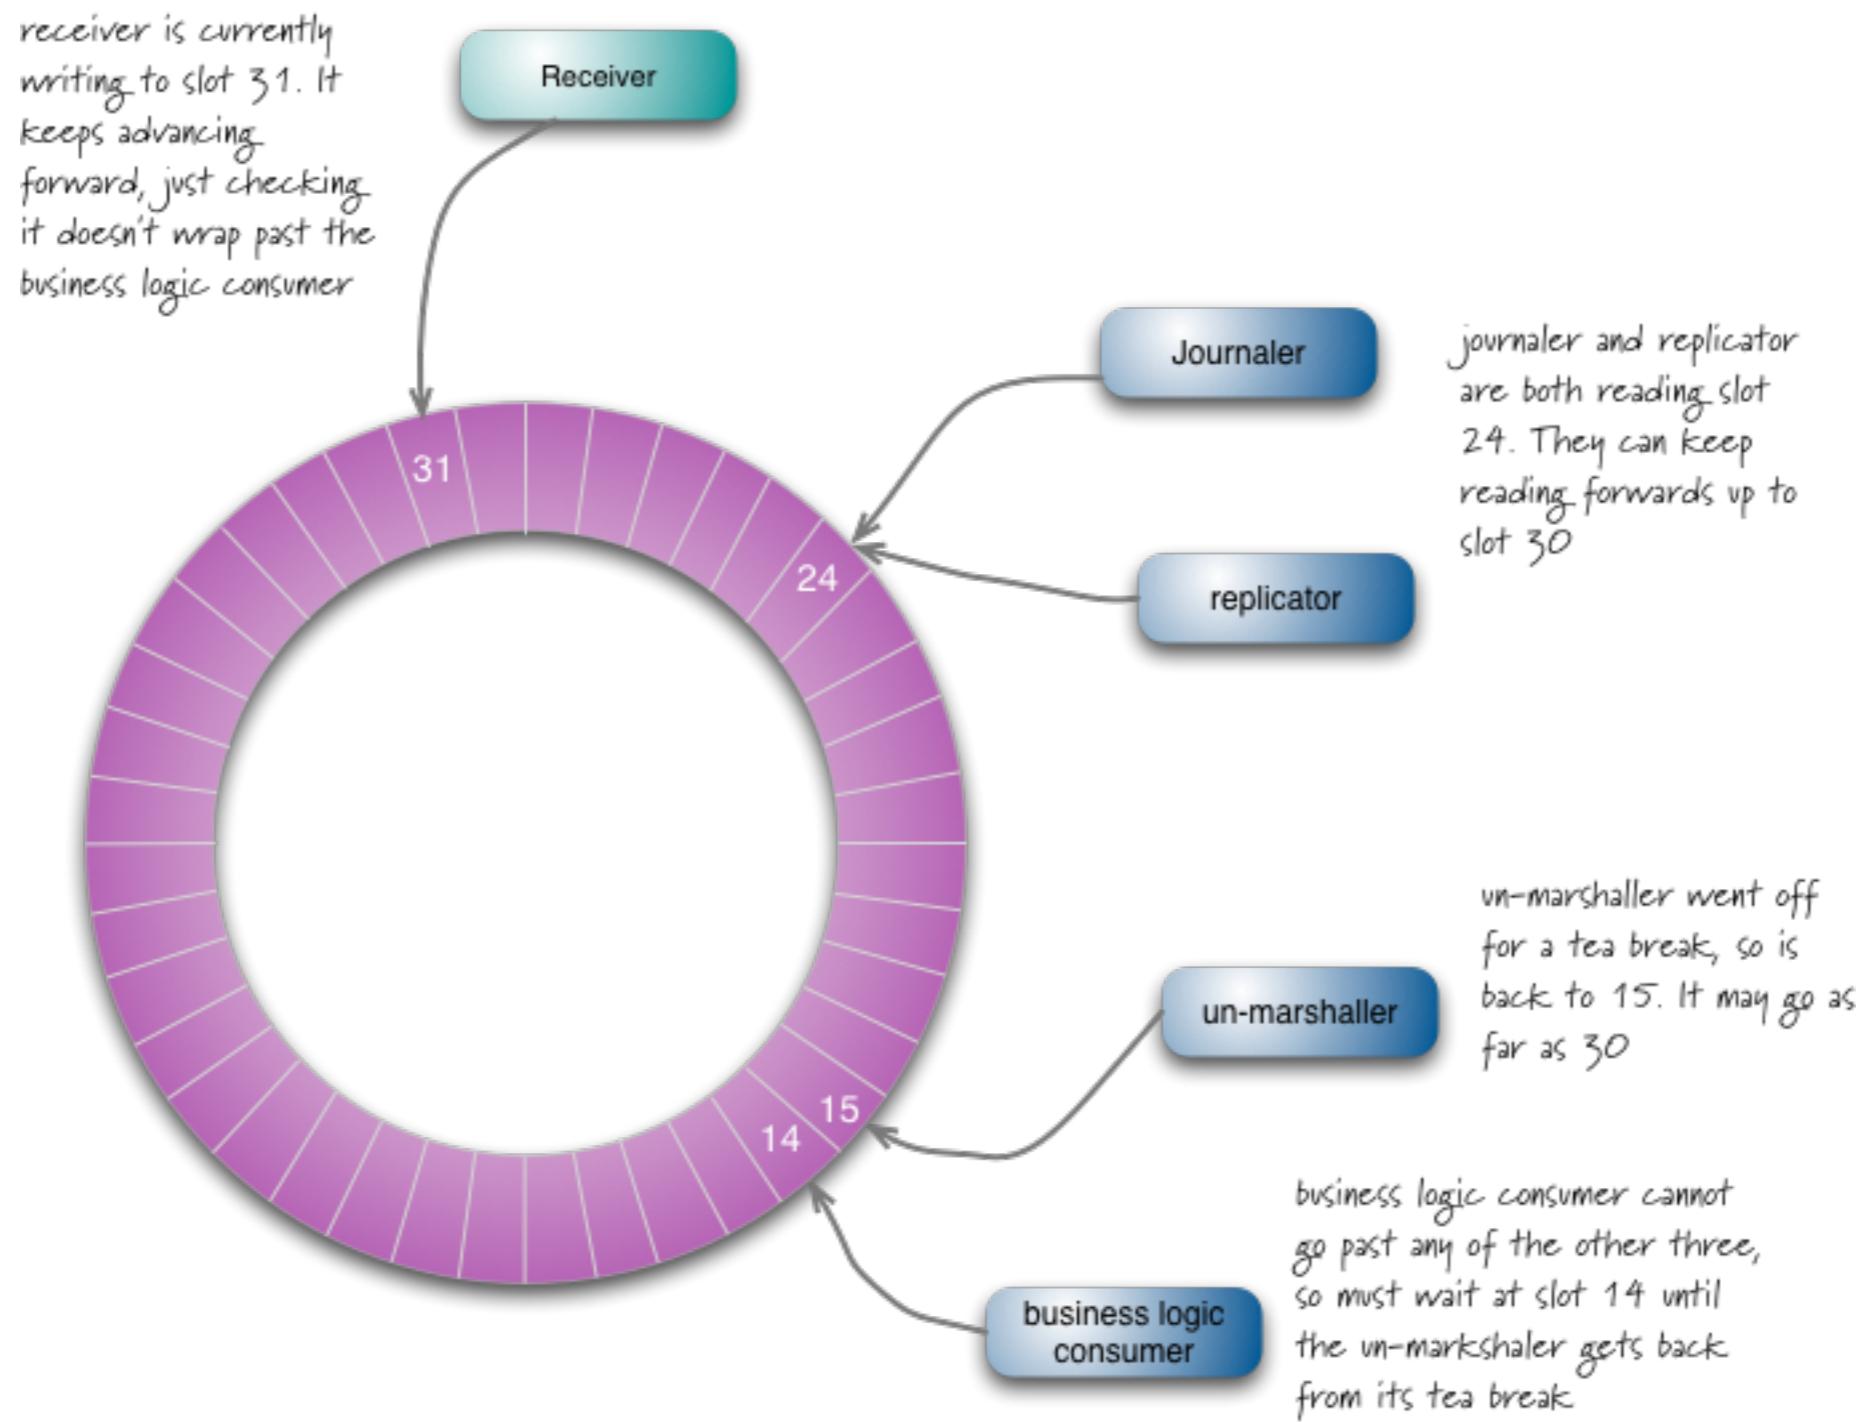
SEDA





Disruptor

Disruptor





Functional Programming

Functional Programming

Referential Transparency
Immutability

Persistent Datastructures



Concurrency models

- Shared state
 - Threads and locks
 - Disruptor
 - FP
 - Message passing
 - SEDA
 - Actors
- WARNING!!! Very incomplete list!!!

Actor Model

Actor Model

- OOP was always about passing messages
 - State encapsulation
- Actors run concurrently
- Actors ***DO*** send messages

What is ?

- fundamental unit of computation
 - process
 - store
 - communicate

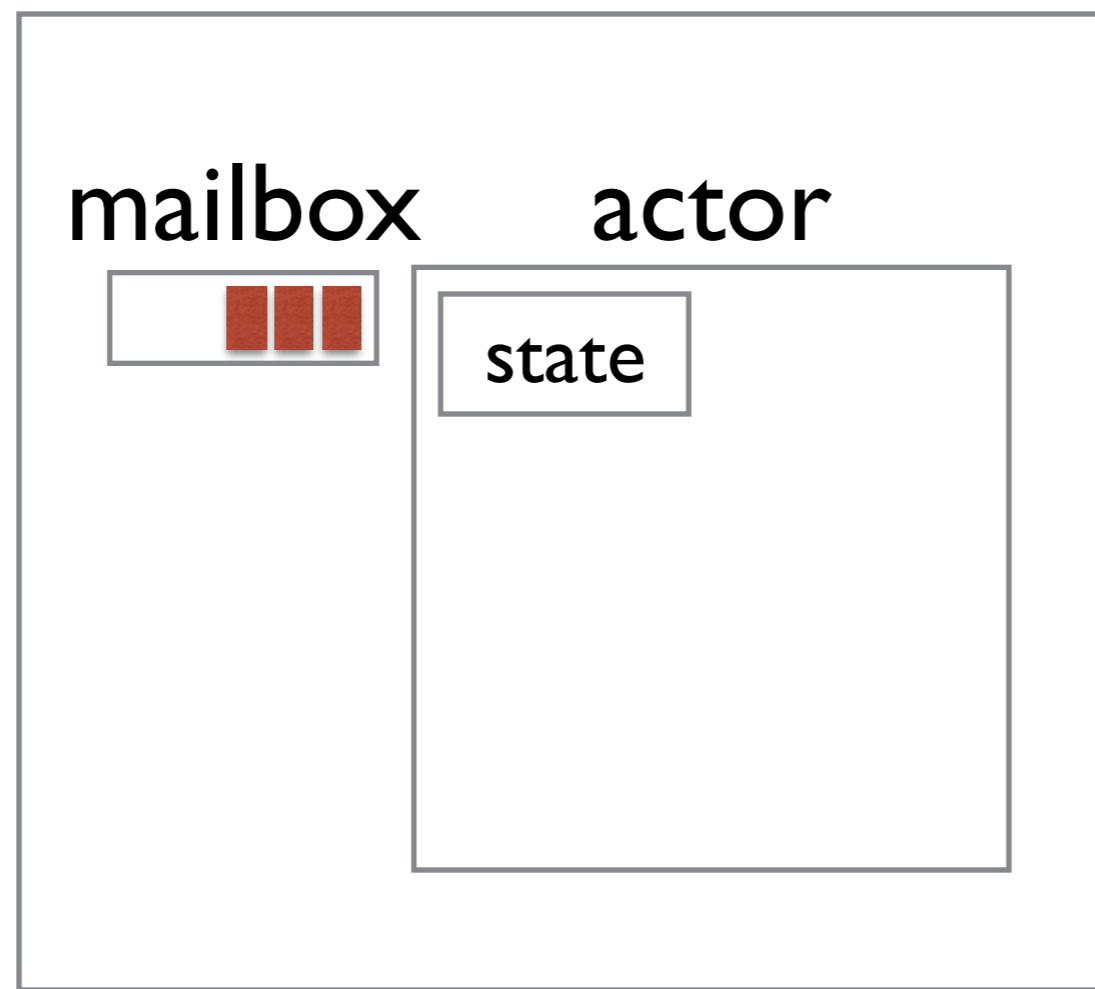
Axioms

- An Actor [...], in response to a message it receives, can concurrently:
 - send messages to addresses of Actors that it has
 - create new Actors
 - designate how to handle the next message it receives

Higher abstraction

- everything runs on a ‘system’
- forget about locks, queues,
- programmer doesn’t know anything about where it’s running. It can be adapted at environment.

address







This is you now
(Only picture without actors)

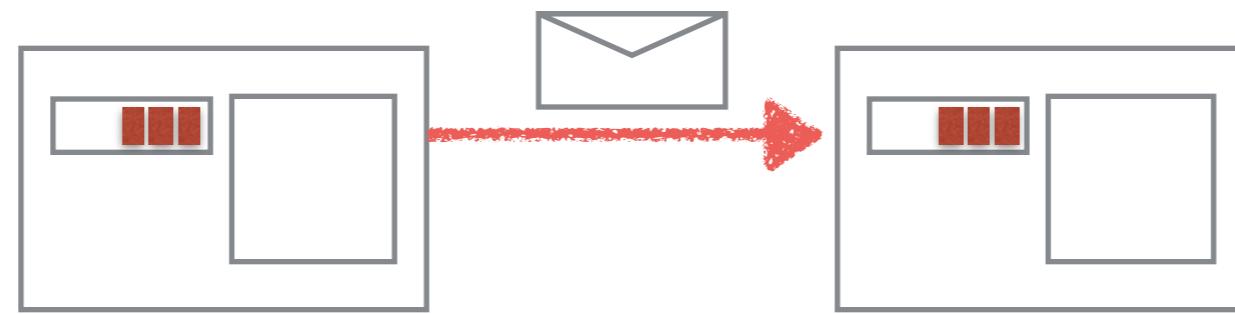


Actors/hello_actors/hello_actors.exs

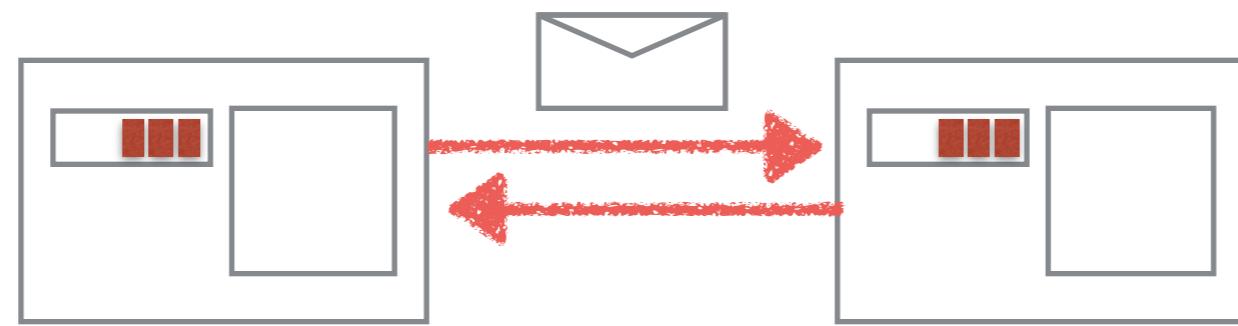
```
def loop do
  receive do
    {:greet, name} -> IO.puts("Hello #{name}")
    {:praise, name} -> IO.puts("#{name}, you're amazing")
    {:celebrate, name, age} -> IO.puts("Here's to another #{age} years, #{name}")
  end
  loop
end
```

source: 7 concurrency
models in 7 weeks

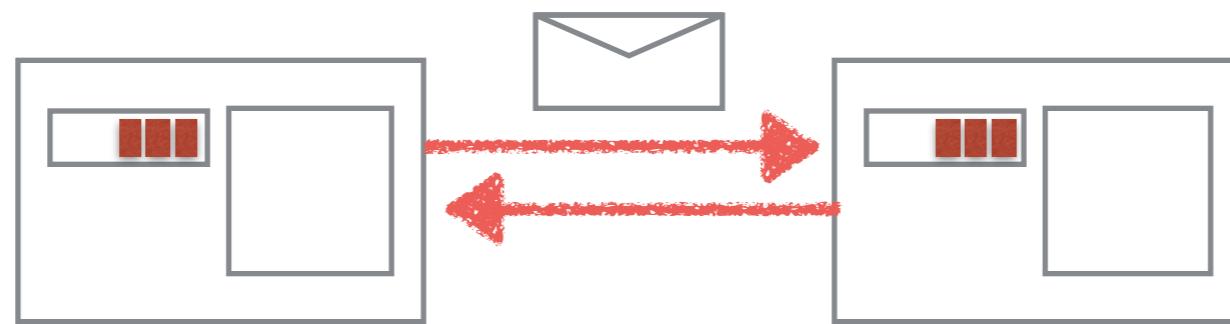
fire and forget



invoke

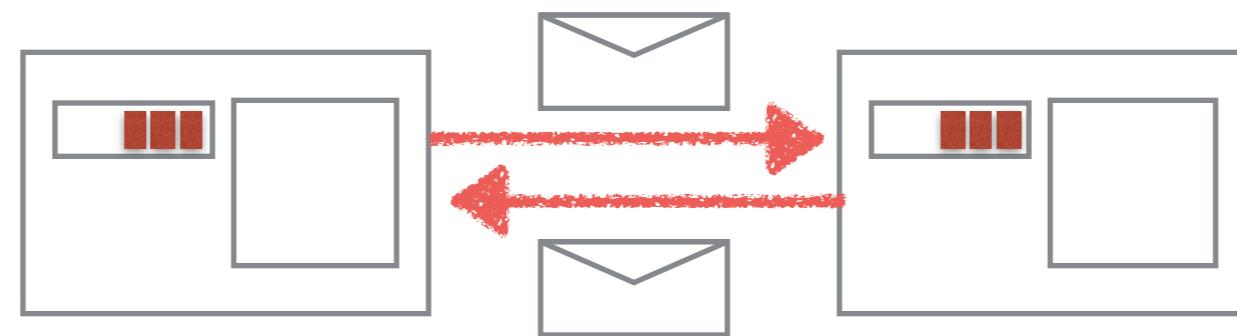


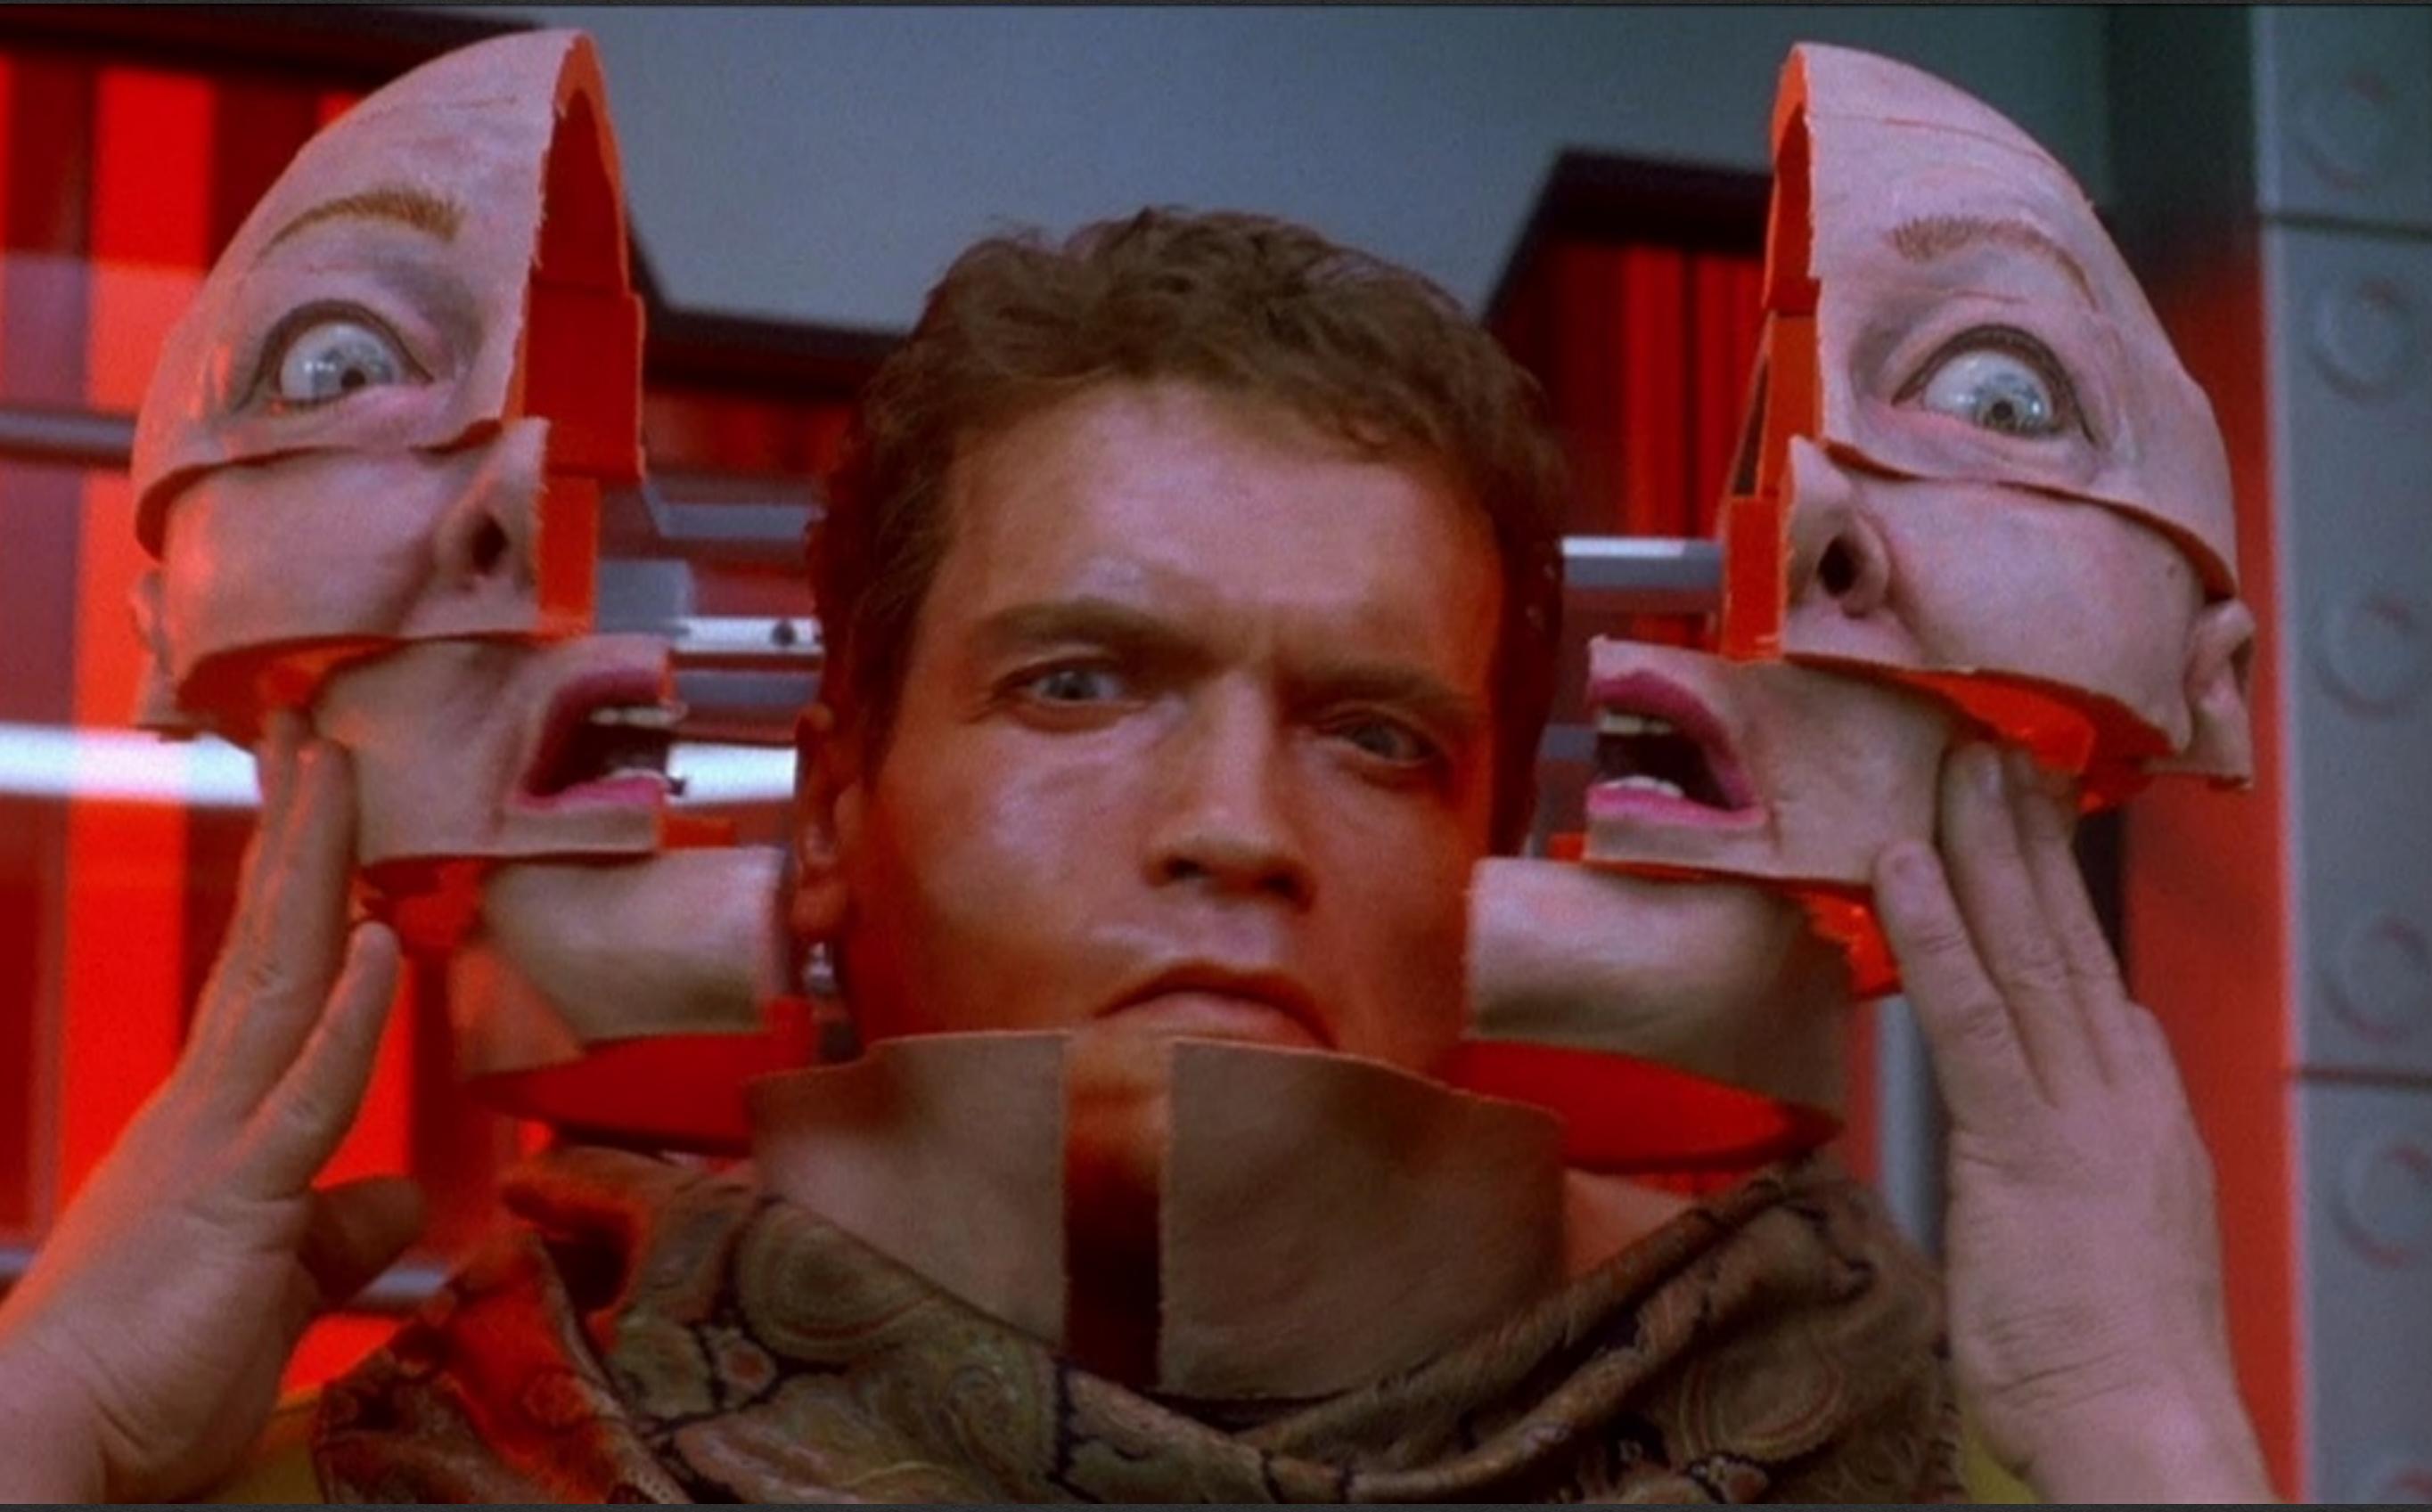
ask



- actors
- futures (for timeouts)
- receiving actor ***must*** reply

ff & ff



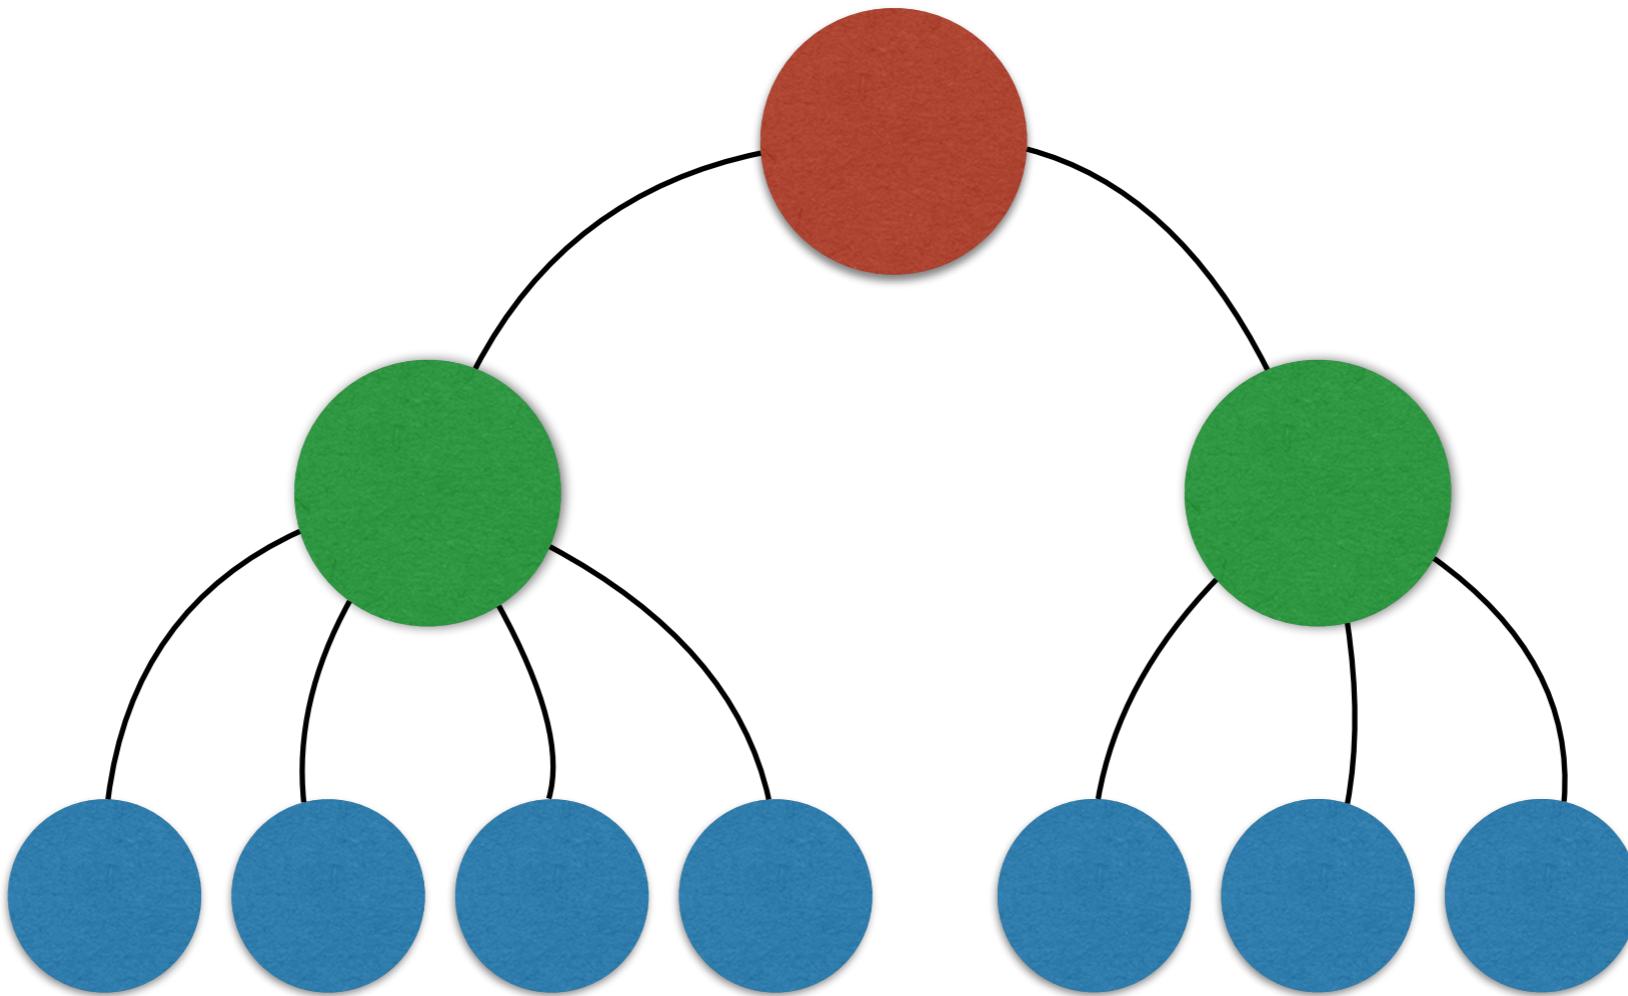


Actors/counter/counter2.ex

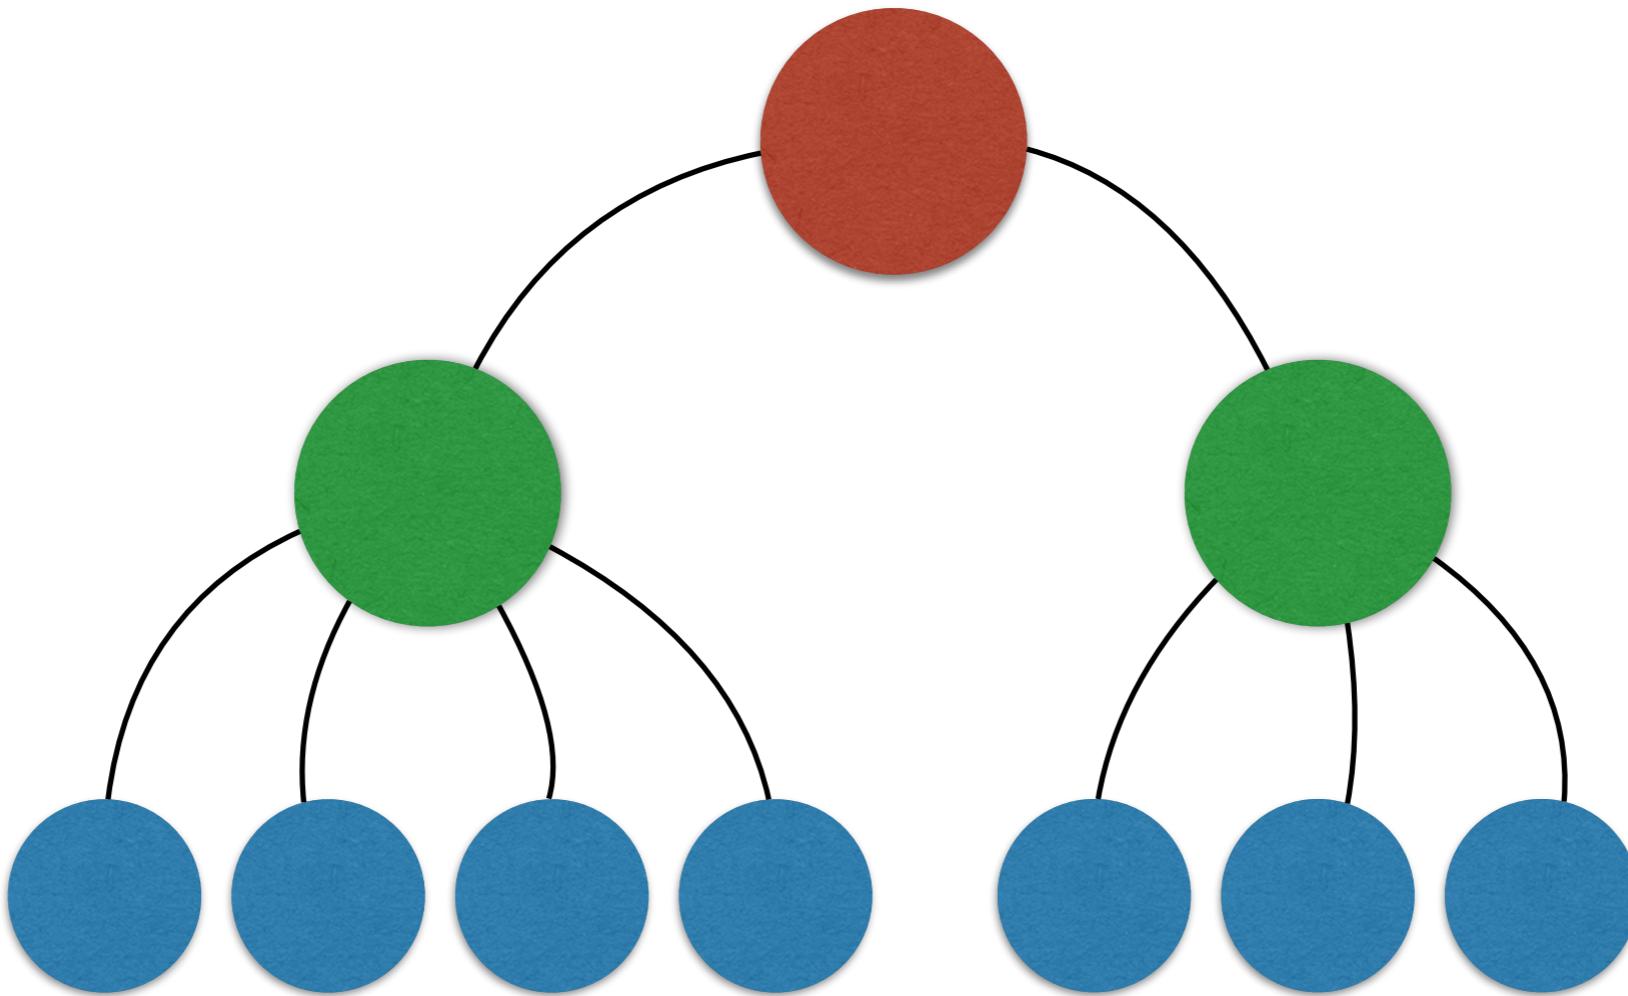
```
defmodule Counter do
  def start(count) do
    spawn(__MODULE__, :loop, [count])
  end
  def next(counter) do
    ▶ ref = make_ref()
    ▶ send(counter, {:next, self(), ref})
    ▶ receive do
      ▶ {:ok, ^ref, count} -> count
    ▶ end
  end
  def loop(count) do
    receive do
      ▶ {:next, sender, ref} ->
        send(sender, {:ok, ref, count})
        loop(count + 1)
    end
  end
end
```

source: 7 concurrency
models in 7 weeks

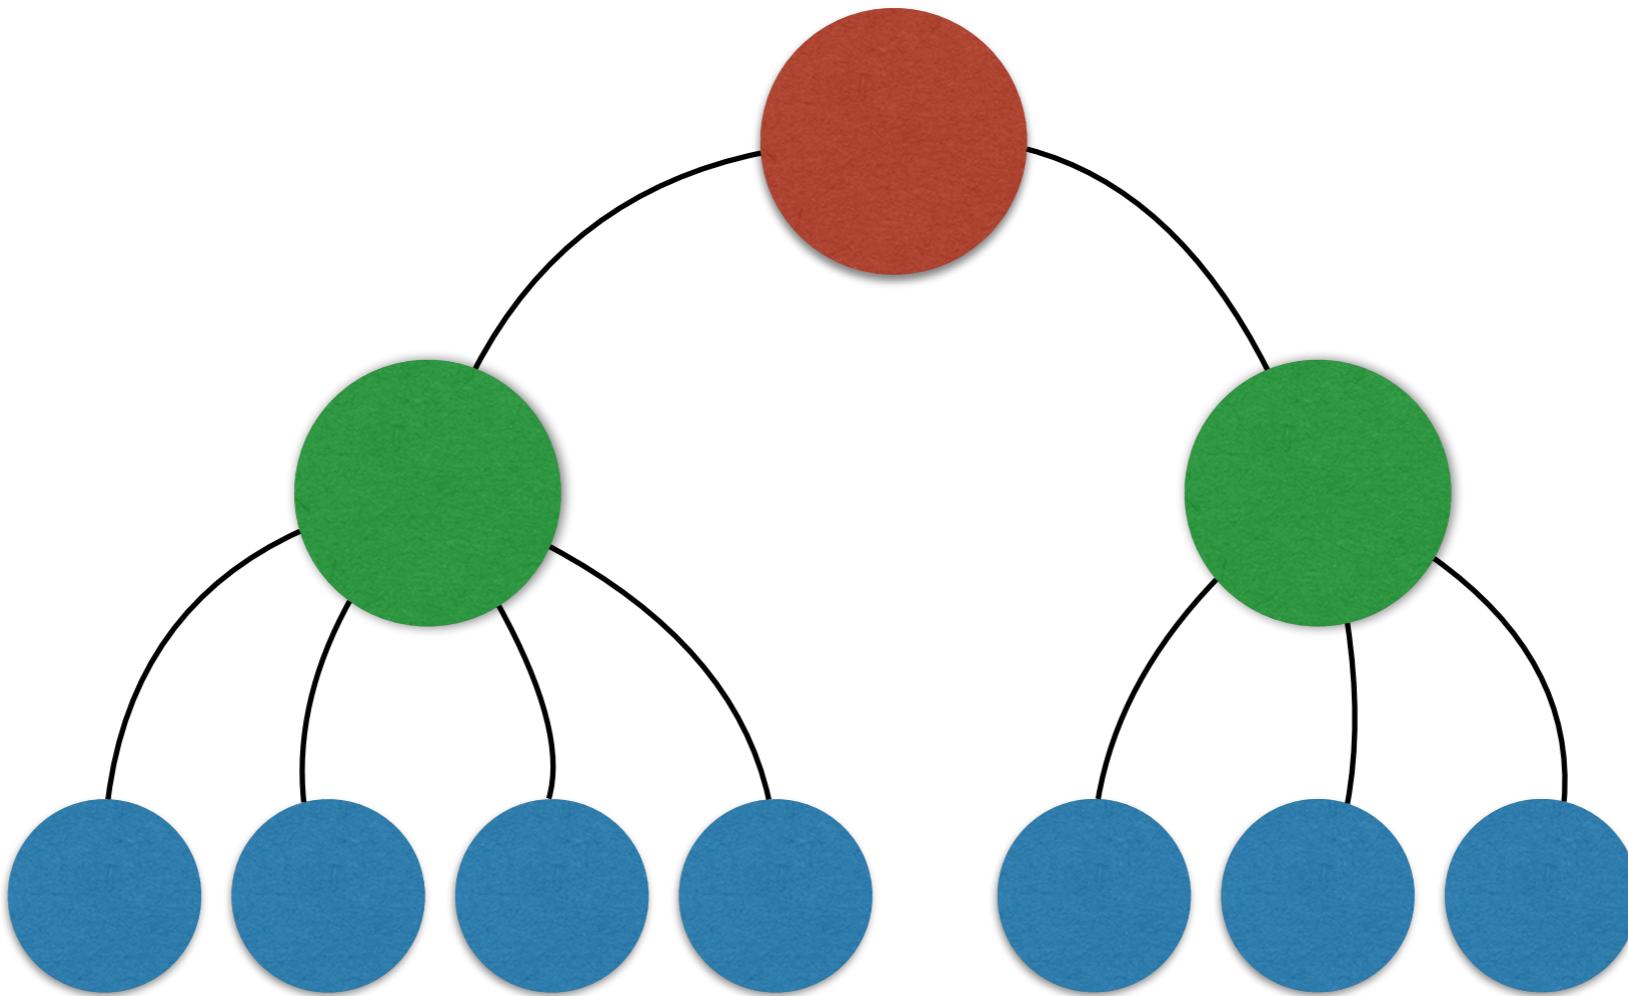
Supervision



One for One



One for One

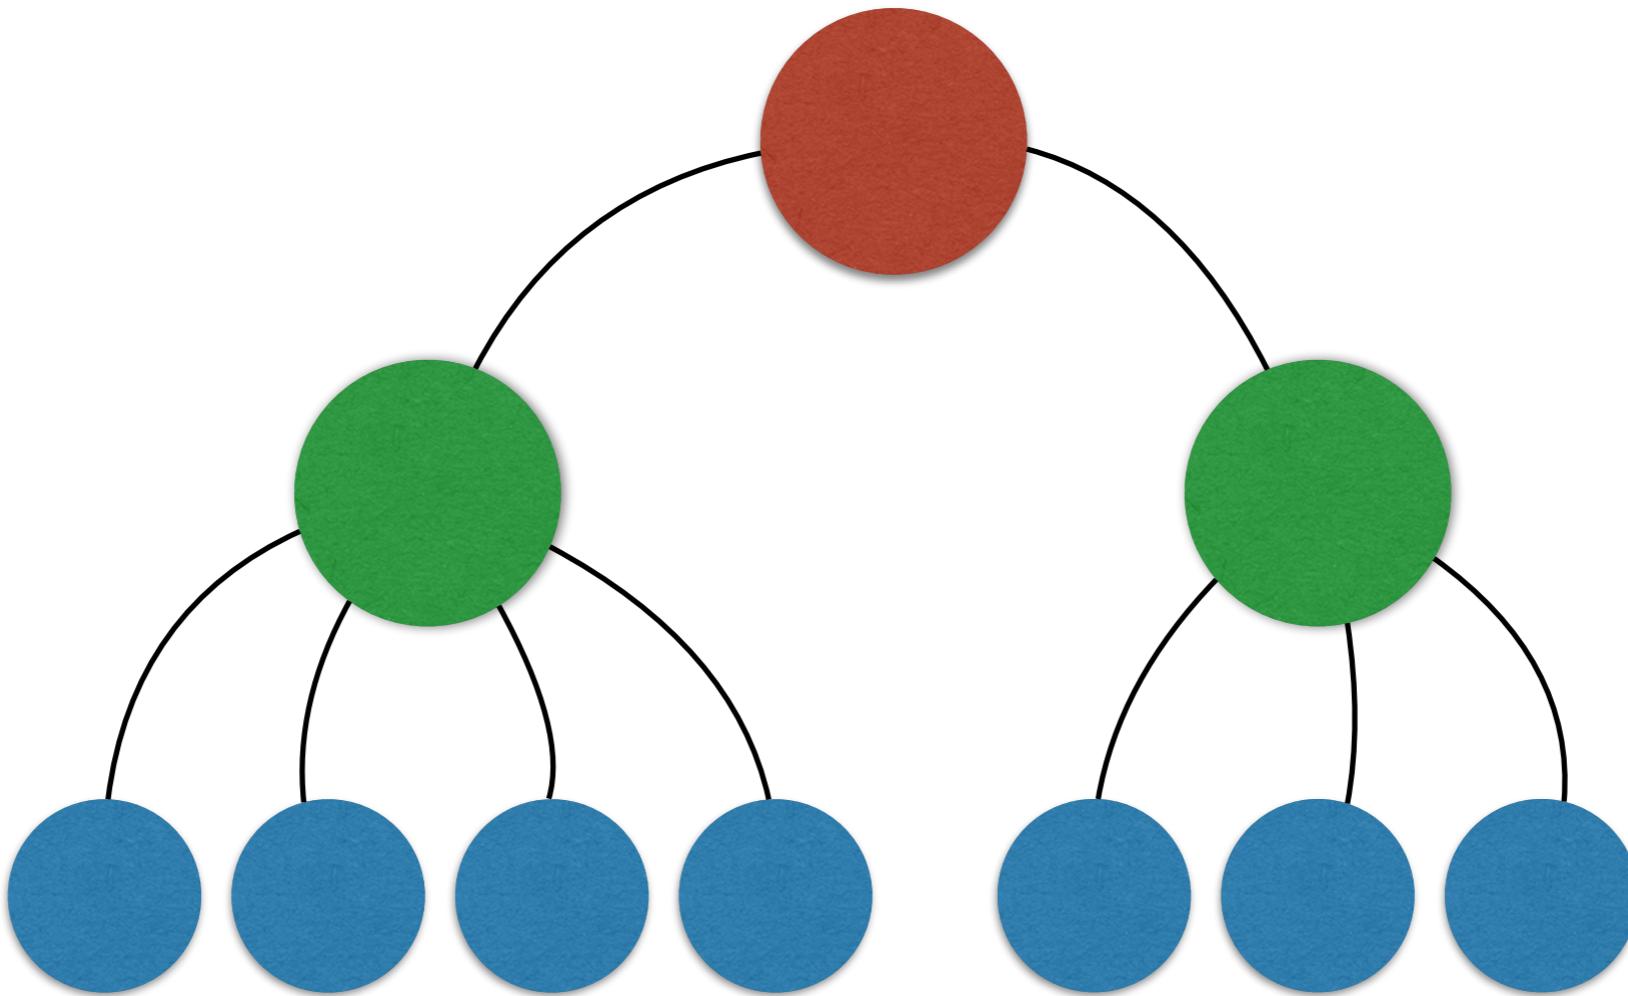


One for One

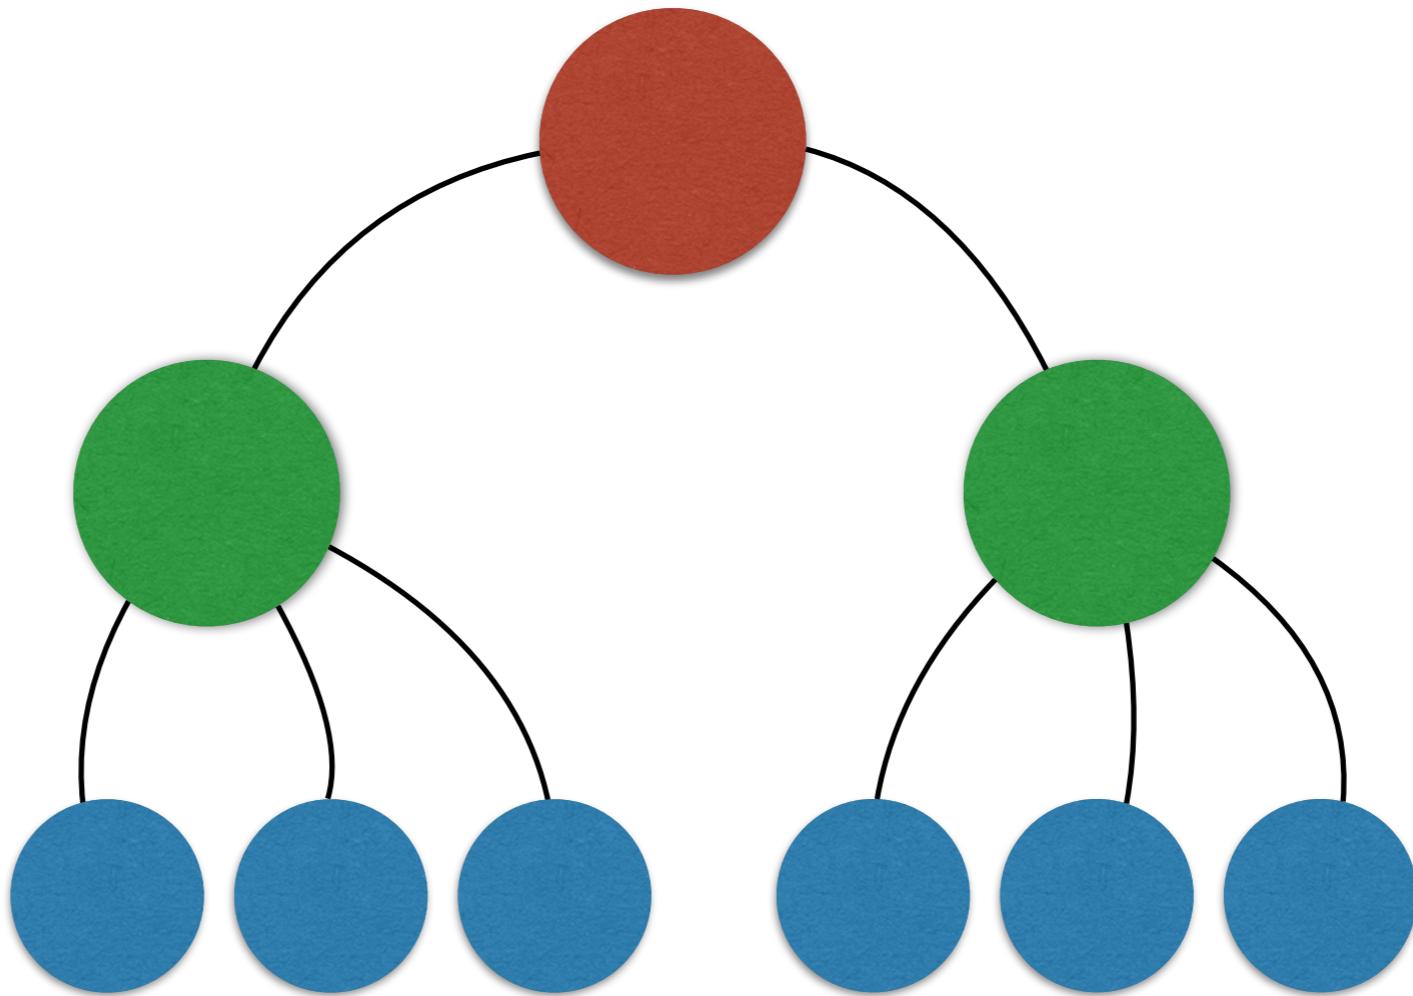
```
4.  
5. override val supervisorStrategy =  
6.   OneForOneStrategy(maxNrOfRetries = 10, withinTimeRange = 1 minute) {  
7.     case _: ArithmeticException      => Resume  
8.     case _: NullPointerException    => Restart  
9.     case _: IllegalArgumentException => Stop  
10.    case _: Exception             => Escalate  
11. }
```

source: akka.io (old version)

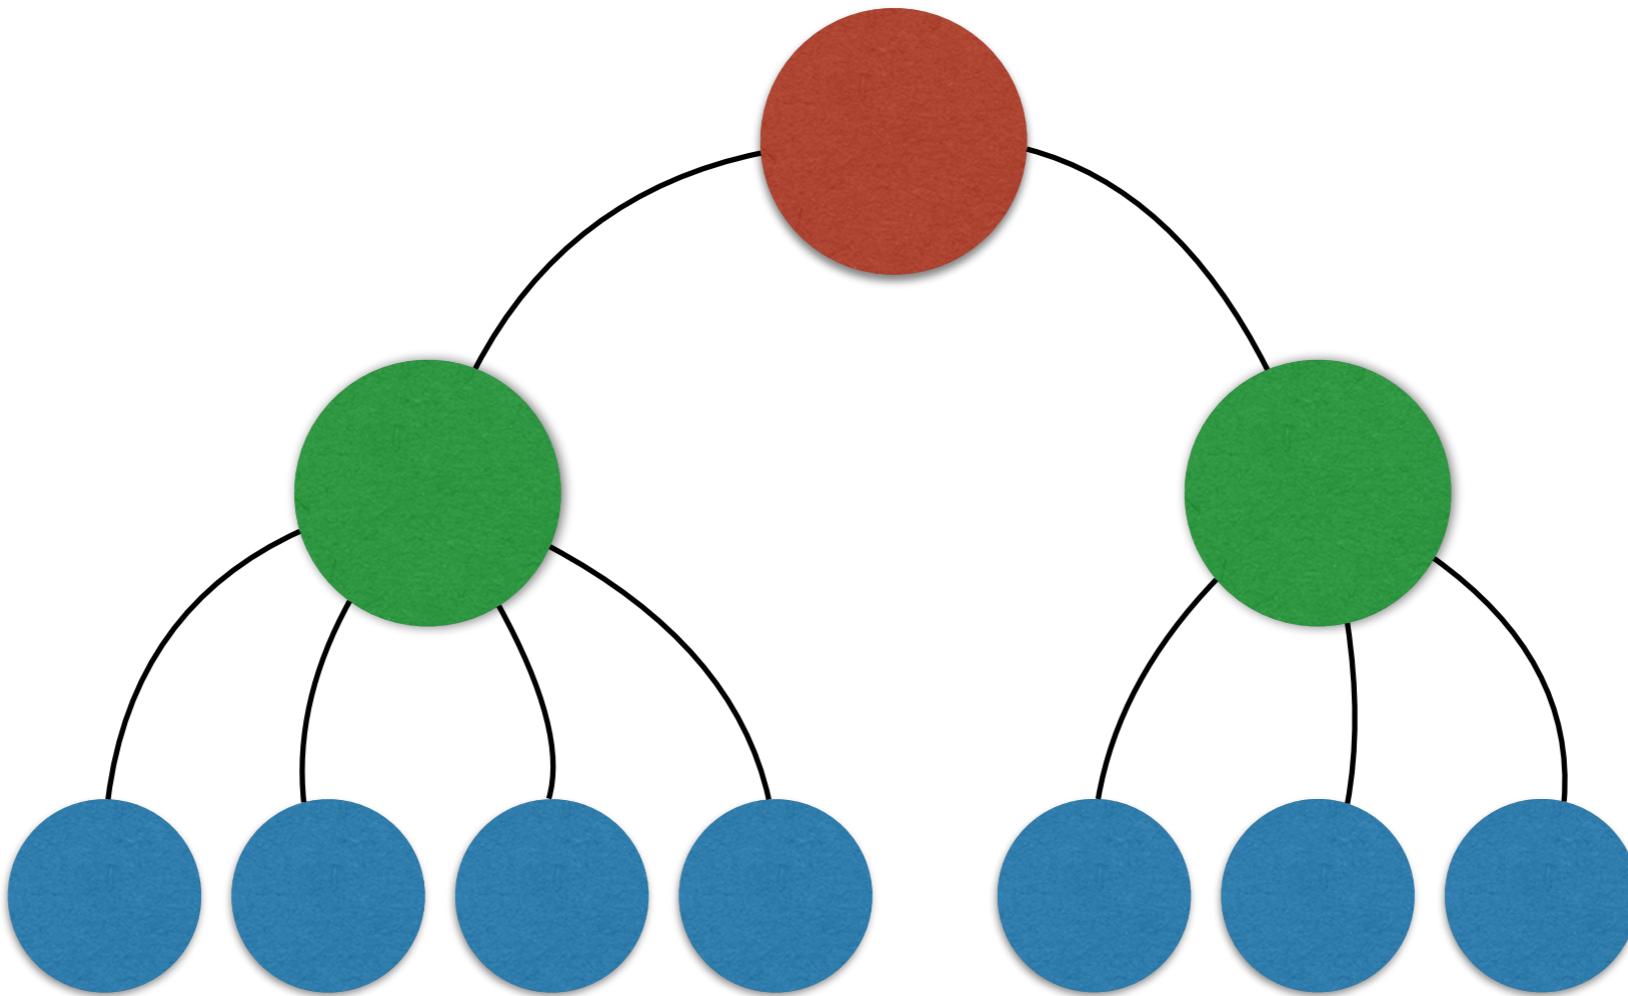
All for One



All for One



All for One



Delivery

at-most-once



Tuning

Dispatcher *
PinnedDispatcher
BalancingDispatcher
CallingThreadDispatcher

Tuning

Dispatcher *

- one MB per actor
- many actors bound to a ThreadPool
- Bulkheading

Tuning

PinnedDispatcher

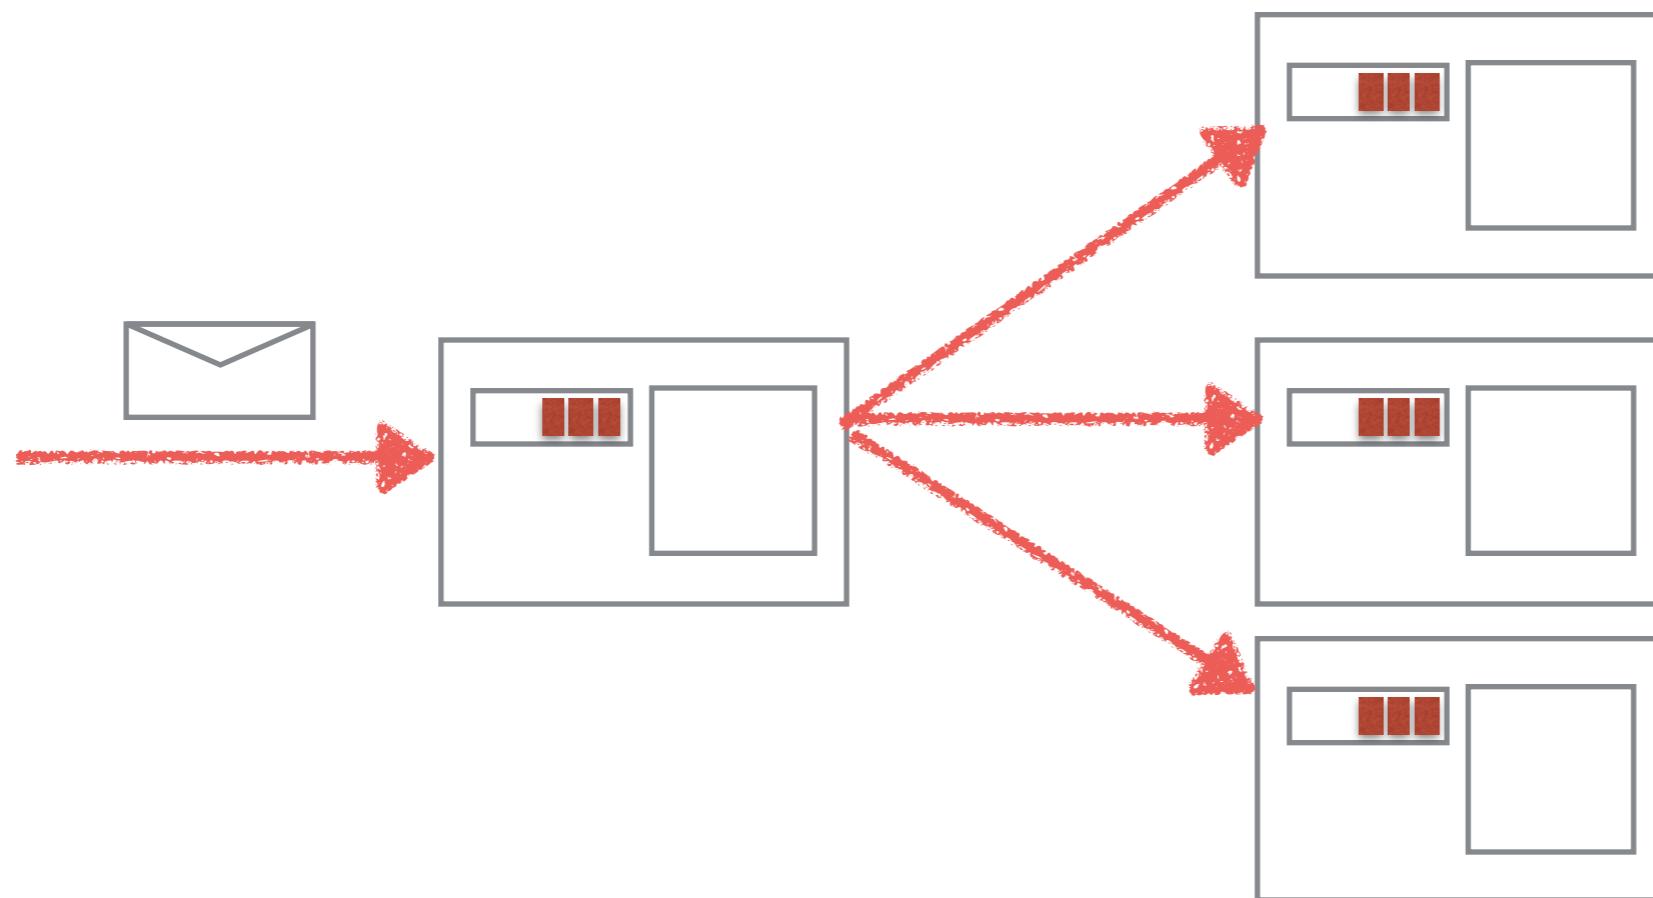
- one MB per actor
- dedicated Thread per actor
- Bulkheading

Tuning

BalancingDispatcher

- one MB - many actors
- ThreadPool
- work-sharing

Patterns



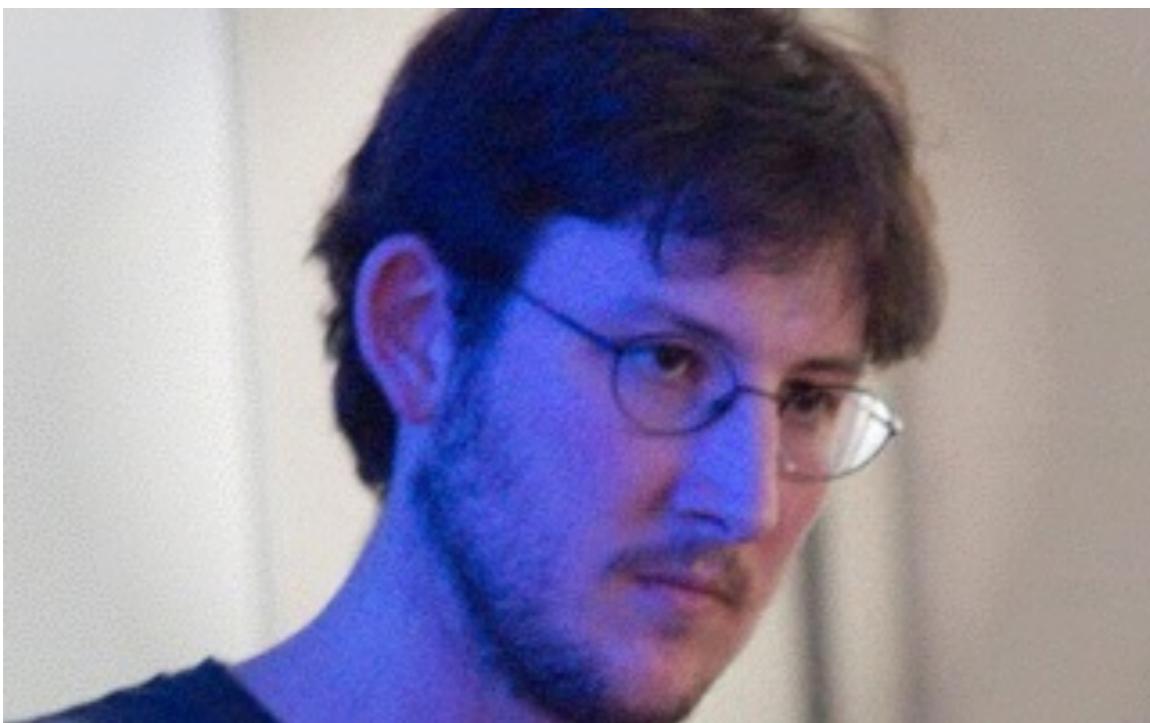
Patterns

- tailchopping
- balancing
- work pulling
- receptionist
- routing
- message forwarding
- ...

Conclusions

- I/O bound vs CPU bound
- YMMV
- know your tools
- actors
 - distributed
 - at-most-once
 - supervision





@ignasi35



Questions?

thanks!

References

<http://docs.oracle.com/cd/E19455-01/806-5257/6je9h032b/index.html>

<https://www.parleys.com/play/51c0bc58e4b0ed877035680a/chapter0/about>

<http://www.slideshare.net/legendofklang/concurrency-scaladays>

<http://dspace.mit.edu/handle/1721.1/6952#files-area>

<http://lmax-exchange.github.io/disruptor/>

<http://martinfowler.com/articles/lmax.html>

http://gotocon.com/dl/goto-chicago-2014/slides/CharlieHunt_TheFundamentalsOfGCPerformance.pdf

<https://www.youtube.com/watch?v=D6jqVS5JFpM&index=7&list=PLEx5khR4g7PJfVDBInUpsY2OG2wlf7hW>